

Approximation Algorithms

Learning Goals.

- Introduce Approximation Algorithms:
 - (PTAS) Polynomial Time Approximation Schemes:
polynomial time algorithms with approximation guarantees
 - Load Balancing (Greedy Balance).
 - Load Balancing (Longest Processing Time First).
 - Vertex Cover (Linear Programming plus Rounding).

Readings: Read Chapter 11 of text, can omit section 11.7.

Based on Chapter 11, Kleinberg and Tardos
 Slides based on those by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved

1

Approximation Algorithms

- Q. Suppose I need to solve an NP-hard problem. What should I do?
A. Theory says you're unlikely to find a poly-time algorithm.

Must sacrifice one of three desired features.

- Solve problem to optimality.
- Solve problem in poly-time.
- Solve arbitrary instances of the problem.

We consider sacrificing optimality, but attempt to preserve a guarantee.

2

Approximation Algorithms (cont.)

ρ -approximation algorithm.

- Guaranteed to run in poly-time.
- Guaranteed to solve arbitrary instance of the problem
- Guaranteed to find solution within ratio ρ of true optimum.

Let C be the computed value and C^* be the optimal, then:

For a maximization problem, $C \geq C^*/\rho$, with $\rho \geq 1$.

For a minimization problem, $C \leq \rho C^*$, with $\rho \geq 1$.

Some authors use $\rho := 1/\rho \leq 1$ for maximization problems.

3

Load Balancing

Input. m identical machines; n jobs, job j has processing time t_j .

- Job j must run contiguously on one machine.
- A machine can process at most one job at a time.

Def. Let $J(i)$ be the subset of jobs assigned to machine i . The

load of machine i is $L_i = \sum_{j \in J(i)} t_j$.

Def. The **makespan** is the maximum load on any machine $L = \max_i L_i$.

Load balancing. Assign each job to a machine to minimize makespan.

4

Load Balancing: List Scheduling

List-scheduling algorithm.

- Consider n jobs in some fixed order.
- Assign job j to machine whose load is smallest so far.



```

List-Scheduling(m, n, t1, t2, ..., tn) {
    for i = 1 to m {
        Li ← 0           ← load on machine i
        J(i) ← ∅          ← jobs assigned to machine i
    }

    for j = 1 to n {
        i = argmink Lk   ← machine i has smallest load
        J(i) ← J(i) ∪ {j}   ← assign job j to machine i
        Li ← Li + tj   ← update load of machine i
    }
    return J, max{Li}
}

```

Implementation. $O(n \log n)$ using a priority queue.

5

Load Balancing: List Scheduling Analysis

Theorem. [Graham, 1966] Greedy algorithm is a 2-approximation.

- First worst-case analysis of an approximation algorithm.
- Need to compare resulting solution with optimal makespan L^* .

Lemma 1. The optimal makespan $L^* \geq \max_j t_j$.

Pf. Some machine must process the most time-consuming job. ■

Lemma 2. The optimal makespan $L^* \geq \frac{1}{m} \sum_j t_j$.

Sketch of Pf.

- The total processing time is $\sum_j t_j$.
- One of m machines must do at least a $1/m$ fraction of total work. ■

These provide useful **lower bounds** on the optimal makespan L^* .


6

Load Balancing: List Scheduling Analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , it had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.



7

Load Balancing: List Scheduling Analysis

Theorem. Greedy algorithm is a 2-approximation.

Pf. Consider load L_i of bottleneck machine i .

- Let j be last job scheduled on machine i .
- When job j assigned to machine i , it had smallest load. Its load before assignment is $L_i - t_j \Rightarrow L_i - t_j \leq L_k$ for all $1 \leq k \leq m$.
- Sum inequalities over all k and divide by m :

$$\begin{aligned}
L_i - t_j &\leq \frac{1}{m} \sum_k L_k \\
&= \frac{1}{m} \sum_i t_i \\
\text{Lemma 2 } \rightarrow &\leq L^*
\end{aligned}$$

- Now $L_i = \underbrace{(L_i - t_j)}_{\leq L^*} + \underbrace{t_j}_{\leq L^*} \leq 2L^*$. ■

↑
Lemma 1

8

Load Balancing: List Scheduling Analysis

Summary: For input s , let $L(s) = L_i$ (the time required for the greedy schedule) and $L^*(s)$ be the minimum finish time. For $\rho = 2$ we have:

$$L(s)/L^*(s) \leq \rho \text{ for all inputs } s.$$

That is, this greedy algorithm is a 2-approximation.

Remaining questions about this algorithm (all roughly similar in nature):

- Can $L(s)$ actually be as bad as $\rho L^*(s)$?
- Is the estimate for ρ sharp?
- Does $\sup\{L(s)/L^*(s) \mid \text{input } s\} = \rho$ (where sup denotes supremum)?
- For a given β in $[1, \rho]$,

$$\boxed{\text{Does there exist an input } s \text{ such that } \beta \leq L(s)/L^*(s)?}$$

E.g., $\beta = 3/2$ (we might try to certify this with one example input s .)

- For a given β in $[1, \rho]$, is $\beta \leq \sup\{L(s)/L^*(s) \mid \text{input } s\}$?


9

Load Balancing: List Scheduling Analysis

Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, first $m(m-1)$ jobs length 1 jobs, then one length m job




10

Load Balancing: List Scheduling Analysis

Q. Is our analysis tight?

A. Essentially yes.

Ex: m machines, $m(m-1)$ jobs length 1 jobs, one job of length m



11

Load Balancing: LPT Rule

Longest processing time (LPT). Sort n jobs in descending order of processing time, and then run list scheduling algorithm.

```

LPT-List-Scheduling(m, n, t1, t2, ..., tn) {
    Sort jobs so that t1 ≥ t2 ≥ ... ≥ tn

    for i = 1 to m {
        Li ← 0           ← load on machine i
        J(i) ← ϕ          ← jobs assigned to machine i
    }

    for j = 1 to n {
        i = argmink Lk      ← machine i has smallest load
        J(i) ← J(i) ∪ {j}       ← assign job j to machine i
        Li ← Li + tj      ← update load of machine i
    }
    return J, max{Li}
}

```

12

Load Balancing: LPT Rule

Observation. If at most m jobs, then list-scheduling is optimal.
Pf. Each job put on its own machine. ■

Lemma 3. If there are more than m jobs, $L^* \geq 2t_{m+1}$.

Pf.

- Consider first $m+1$ jobs t_1, \dots, t_{m+1} .
- Since the t_i 's are in descending order, each takes at least t_{m+1} time.
- There are $m+1$ jobs and m machines, so by pigeonhole principle, at least one machine gets two jobs. ■

Theorem. LPT rule is a $3/2$ approximation algorithm.

Pf. Same basic approach as for list scheduling.

$$L_i = \underbrace{(L_i - t_j)}_{\leq L^*} + \underbrace{t_j}_{\leq \frac{3}{2}L^*} \leq \frac{3}{2}L^*. \quad \blacksquare$$

↑
Lemma 3
(by observation, can assume number of jobs > m)

13

Load Balancing: LPT Rule

Q. Is our $3/2$ analysis tight? I.e., is $\sup\{L(s)/L^*(s) \mid s\} = 3/2$?
A. No.

Theorem. [Graham, 1969] LPT rule is a $4/3$ -approximation.

Pf. More sophisticated analysis of same algorithm.

Q. Is Graham's $4/3$ analysis tight?


A. Yes.

Ex: m machines, $n = 2m$ jobs, with one job of length $3m$, another job of length m , and then 2 jobs of each length $m+1, m+2, \dots, 2m-1$.

14

Approximating Vertex Cover

Vertex Cover. Given an undirected graph $G = (V, E)$, find a minimum size subset of nodes S such that every edge is incident to at least one vertex in S .




Minimum size vertex cover?

$$S^* = \{ A, C, G, J \}$$

15

Greedy by Degree Algorithm

Greedy by Degree Algorithm for Vertex Cover: Select next vertex to be one with maximum degree. Remove it and its corresponding edges. Repeat until the remaining graph has no edges.




Step 1: $S = \{ G \}$

16

Greedy by Degree Algorithm

Greedy by Degree Algorithm for Vertex Cover: Select next vertex to be one with maximum degree. Remove it and its corresponding edges. Repeat until the remaining graph has no edges.




Step 2: $S = \{ G, C \}$

17

Greedy by Degree Algorithm

Greedy by Degree Algorithm for Vertex Cover: Select next vertex to be one with maximum degree. Remove it and its corresponding edges. Repeat until the remaining graph has no edges.




Step 3: $S = \{ G, C, J \}$

18

Greedy by Degree Algorithm

Greedy by Degree Algorithm for Vertex Cover: Select next vertex to be one with maximum degree. Remove it and its corresponding edges. Repeat until the remaining graph has no edges.



Step 4: $S = \{ G, C, J, A \}$

19

Greedy by Degree Algorithm

Greedy by Degree Algorithm for Vertex Cover: Select next vertex to be one with maximum degree. Remove it and its corresponding edges. Repeat until the remaining graph has no edges.



Return: $S = \{ G, C, J, A \}$

Same as $S^* = \{ A, C, G, J \}$ for this example

20

Analysis of Greedy by Degree

Is the Greedy by Degree algorithm a ρ -approximation?

That is, for

- Input s describing a graph $G = (V, E)$.
- $L(s)$ defined to be $|S|$, where S is the output of greedy by degree.
- $L^*(S)$ defined to be $|S^*|$, where S^* is a minimum sized vertex cover.

Then, does there exist a constant ρ such that

$$L(s)/L^*(s) \leq \rho \text{ for all inputs } s?$$

Answer: No. Blackboard: Graphs $s=(V,E)$ s.t. $L(s)/L^*(s) \in \Theta(\log|V|)$.


Conclusion: The Greedy by Degree algorithm is not a ρ -approximation, for any constant ρ . Turns out it is a $O(\log(|V|))$ -approximation.

21

Edge Removal Algorithm

Edge Removal Algorithm for Vertex Cover: Select an edge. Add both endpoints to vertex cover. Remove all edges terminating at these endpoints . Repeat until the remaining graph has no edges.

(E.g. Select edges in lexical order: ((A,F), (A,G), (B,C), ... (I, J))




Step 1: $S = \{ A, F \}$

22

Edge Removal Algorithm

Edge Removal Algorithm for Vertex Cover: Select an edge. Add both endpoints to vertex cover. Remove all edges terminating at these endpoints . Repeat until the remaining graph has no edges.

(E.g. Select edges in lexical order: ((A,F), (A,G), (B,C), ... (I, J))




Step 2: $S = \{ A, F, B, C \}$

23

Edge Removal Algorithm

Edge Removal Algorithm for Vertex Cover: Select an edge. Add both endpoints to vertex cover. Remove all edges terminating at these endpoints . Repeat until the remaining graph has no edges.

(E.g. Select edges in lexical order: ((A,F), (A,G), (B,C), ... (I, J))




Step 3: $S = \{ A, F, B, C, D, G \}$

24

Edge Removal Algorithm

Edge Removal Algorithm for Vertex Cover: Select an edge. Add both endpoints to vertex cover. Remove all edges terminating at these endpoints . Repeat until the remaining graph has no edges.
 (E.g. Select edges in lexical order: ((A,F), (A,G), (B,C), ... (I, J))




Step 4: $S = \{ A, F, B, C, D, G, E, J \}$

25

Edge Removal Algorithm

Edge Removal Algorithm for Vertex Cover: Select an edge. Add both endpoints to vertex cover. Remove all edges terminating at these endpoints . Repeat until the remaining graph has no edges.
 (E.g. Select edges in lexical order: ((A,F), (A,G), (B,C), ... (I, J))



Apply this alg to the examples used to show Greedy by Degree is not a p-approximation

Return: $S = \{ A, F, B, C, D, G, E, J \}, |S| = 8$.

Recall $S^* = \{ A, C, G, J \}$ for this example, $|S^*| = 4$.

26

Edge Removal is a 2-approximation

Theorem. If S^* is a vertex cover of minimum size, then any S produced by the Edge Removal Alg. is a vertex cover with $|S| \leq 2|S^*|$.

Pf. $[S$ is a vertex cover]

- In the alg, an edge $(u, v) \in E$ is removed only when covered.
- Algorithm terminates when there are no edges left, with all edges covered.

Pf. $[|S| \text{ is at most size } 2|S^*|]$

- Let S^* be optimal vertex cover.
- Let E_S be the set of edges selected by Edge Removal (where both endpoints are added to S).
- For every $e \in E_S$, then e must have at least one endpoint in S^* :

$$|S| = \sum_{e \in E_S} 2 = 2|E_S| \leq 2|S^*|$$

Both endpoints of $e \in E_S$ in S


Each $e \in E_S$ has at least one endpoint in S^*

Moral: Less greed is sometimes better (e.g., this vs greedy-by-degree).

27

Weighted Vertex Cover

Weighted vertex cover. Given an undirected graph $G = (V, E)$ with vertex i having weight $w_i \geq 0$, find a minimum weight subset of nodes S such that every edge is incident to at least one vertex in S .



Total weight $(9 + 10 + 6 + 32) = 57$. Minimum?

28

Weighted Vertex Cover: ILP Formulation

Weighted vertex cover. Given an undirected graph $G = (V, E)$ with vertex i having weight $w_i \geq 0$, find a minimum weight subset of nodes S such that every edge is incident to at least one vertex in S .

Integer linear programming formulation.

- Model inclusion of each vertex i using a 0/1 variable x_i .

$$x_i = \begin{cases} 0 & \text{if vertex } i \text{ is not in vertex cover} \\ 1 & \text{if vertex } i \text{ is in vertex cover} \end{cases}$$

Vertex covers in 1-1 correspondence with 0/1 assignments:

$$S = \{i \in V : x_i = 1\}$$

- Objective function: minimize $\sum_i w_i x_i$.
- For each edge $(i, j) \in E$, must take either i or j : $x_i + x_j \geq 1$.

29

Weighted Vertex Cover: ILP Formulation

Weighted vertex cover. Integer linear programming formulation.

$$\begin{aligned} (\text{ILP}) \quad \min \quad & \sum_{i \in V} w_i x_i \\ \text{s. t.} \quad & x_i + x_j \geq 1 \quad (i, j) \in E \\ & x_i \in \{0, 1\} \quad i \in V \end{aligned}$$

Observation. If x^* is optimal solution to (ILP), then $S = \{i \in V : x^*_i = 1\}$ is a min weight vertex cover.

30

Weighted Vertex Cover: Real-Valued Relaxation

Weighted vertex cover. Linear programming formulation.


$$\begin{aligned} (\text{LP}) \quad \min \quad & \sum_{i \in V} w_i x_i \\ \text{s. t.} \quad & x_i + x_j \geq 1 \quad (i, j) \in E \\ & x_i \geq 0 \quad i \in V \end{aligned}$$

Note, $x_i \leq 1$ is not necessary.

Observation. Optimal value of (LP) is \leq optimal value of (ILP).

Pf. LP has fewer constraints.

Note. LP is not equivalent to vertex cover.



Q. How can solving LP help us find a small vertex cover?

A. Solve LP and round fractional values.

31

LP + Rounding is a 2-Approx for Weighted Vertex Cover

Theorem. If x^* is optimal solution to (LP), then $S = \{i \in V : x^*_i \geq \frac{1}{2}\}$ is a vertex cover whose weight W_{RLP} is at most twice the min possible weight W^* .

Pf. [S is a vertex cover]

- Consider an edge $(i, j) \in E$.
- Since $x^*_i + x^*_j \geq 1$, either $x^*_i \geq \frac{1}{2}$ or $x^*_j \geq \frac{1}{2} \Rightarrow (i, j)$ covered.

Pf. [S has desired cost]

- Let S^* be optimal vertex cover. Then

$$W^* = \sum_{i \in S^*} w_i \geq \sum_{i \in V} w_i x_i^* \geq \sum_{i \in S} w_i x_i^* \geq \frac{1}{2} \sum_{i \in S} w_i = \frac{1}{2} W_{RLP}$$

LP is a relaxation

$x^* \geq \frac{1}{2}$

- Therefore $W_{RLP} \leq 2W^*$ and the rounded LP algorithm is a 2-approximation.

32

Weighted Vertex Cover

Theorem (above). There exists a 2-approximation algorithm for weighted vertex cover.

Theorem. [Dinur-Safra 2001] If $P \neq NP$, then no ρ -approximation for $\rho < 1.3607$, even with unit weights.

10 √5 - 21

Open research problem. Close the gap.


33

Extra Slides

Pricing Method for Weighted Vertex Cover

Pricing method. Each edge must be covered by some vertex i . Edge e pays price $p_e \geq 0$ to use vertex i .

Fairness. Edges incident to vertex i should pay $\leq w_i$ in total.



Claim. For any vertex cover S and any fair prices p_e : $\sum_e p_e \leq w(S)$.

Proof.

$$\sum_{e \in E} p_e \leq \sum_{i \in S} \sum_{e=(i,j)} p_e \leq \sum_{i \in S} w_i = w(S).$$

↑ ↑
each edge e covered by sum fairness inequalities
at least one node in S for each node in S

35

Pricing Method

Pricing method. Set prices and find vertex cover simultaneously.

```
Weighted-Vertex-Cover-Approx(G, w) {
    foreach e in E
        p_e = 0
    while (∃ edge i-j such that neither i nor j are tight)
        select such an edge e
        increase p_e without violating fairness
    S ← set of all tight nodes
    return S
}
```

$\sum_{e=(i,j)} p_e = w_i$

36

Pricing Method




Figure 11.8

37

Pricing Method: Analysis

Theorem. Pricing method is a 2-approximation.

Pf.

- Algorithm terminates since at least one new node becomes tight after each iteration of while loop.
- Let S = set of all tight nodes upon termination of algorithm. S is a vertex cover: if some edge $i-j$ is uncovered, then neither i nor j is tight. But then while loop would not terminate.
- Let S^* be optimal vertex cover. We show $w(S) \leq 2w(S^*)$.

$$w(S) = \sum_{i \in S} w_i = \sum_{i \in S} \sum_{e=(i,j)} p_e \leq \sum_{i \in V} \sum_{e=(i,j)} p_e = 2 \sum_{e \in E} p_e \leq 2w(S^*) \quad \blacksquare$$

↓
 all nodes in S are tight
 ↓
 $S \subseteq V$,
 prices ≥ 0
 ↑
 each edge counted twice
 ↑
 fairness lemma

38