

Local Search

Learning Goals.

- Introduce Local Search.
- Metropolis Algorithm and Simulated Annealing.
- Max-Cut Problem.
- Local Search in Shape from Texture.

Readings: Reading Chp. 12, up to end of Section 12.5.

Based on Chapter 12, Kleinberg and Tardos
PEARSON Addison Wesley
Slides based on those by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved

1

Coping With NP-Hardness

- Q. Suppose I need to solve an NP-hard problem. What should I do?
A. Theory says you're unlikely to find poly-time algorithm.

Must sacrifice one of three desired features.

- Solve problem to optimality.
- Solve problem in polynomial time.
- Solve arbitrary instances of the problem.

Here we use [iterative local search](#), finding a slightly better solution each step, until we cannot locally improve the solution.

Blind hill-climbing (or descent), usually with weak or no guarantees.

A key component here is how the [local search neighbourhood](#) is defined.


2

Local Search and Local Minima

Local search. Algorithm that explores the space of possible solutions in sequential fashion, moving from a current solution to a "nearby" one.

Neighbour relation. Let $S \sim S'$ be a neighbor relation for the problem. E.g., a unit-length left/right move for the ball below.

Greedy descent. Let S denote current solution. If there is a neighbor S' of S with strictly lower cost, replace S with the neighbor whose cost is as small as possible. Otherwise, terminate the algorithm.



3

Simplex Method as Local Search

Maximize LP by switching between neighbouring vertices.


Problem Constraints:
 $x_1 + 3x_3 \leq 600$
 $x_2 + x_3 \leq 300$
 $x_1 + x_2 + x_3 \leq 400$
 $x_2 \leq 250$

Feasible set F is this closed polytope.

Maximizes $c^T x$,
 $c^T = (2, 3, 4)$,
 $c^T x$ gives shading.

Convexity of problem guarantees convergence to optimal solution.

Convex (and "pseudo"-convex) problems are like funnels.




4

Greedy Descent: Vertex Cover

Vertex Cover Problem: Given a graph $G = (V, E)$, find a subset of nodes S of minimal cardinality such that for each $(u, v) \in E$, either u or v (or both) are in S .

Example: S initially V .



Neighbour relation. Vertex cover $S \sim S'$ if S' can be obtained from S by adding or deleting a single vertex and S' itself is a vertex cover. Each vertex cover S has at most $n-1$ neighbors (as usual, $|V| = n$).

5

Greedy Descent: Vertex Cover

Vertex Cover Greedy Algorithm:

```
Greedy_Vertex_Cover(V, E) {
    S ← V
    while (true) {
        Nhbrs = {S' s.t. S' ~ S, |S'| < |S| and
                 S' is a vertex cover}.
        if Nhbrs is not empty
            S ← an element of Nhbrs
        else
            return S
    }
}
```


Remark. Algorithm terminates after at most n steps since each update decreases the size of the cover by one.

Key Question: How close to optimal is the resulting vertex cover?

6

Greedy Descent: Vertex Cover

Example: S initially V .




Remark. Algorithm can get stuck in a "local minimum". Might consider a larger neighbourhood to try to avoid this. E.g., Swap up to two vertices into or out of S . Then greedy descent might need to check $O(|V|^2)$ pairs.

Might consider randomly selecting amongst ties. Running several times.


7

Greedy Descent: Vertex Cover


Local optimum. No neighbour (a cover) is strictly better.



optimum = center node only
local optimum = all other nodes



optimum = all nodes on left side
local optimum = all nodes on right side



optimum = even nodes $|S| = 4$
local optimum = odd nodes $|S| = 5$

8

Metropolis Algorithm

Metropolis algorithm. [Metropolis, Rosenbluth, Rosenbluth, Teller, Teller 1953]

- Simulate behavior of a physical system according to principles of statistical mechanics.
- Globally biased toward "downhill" steps, but occasionally makes "uphill" steps to break out of local minima.

Gibbs-Boltzmann function. The probability of finding a physical system in a state with energy E is proportional to $e^{-E/(kT)}$, where $T > 0$ is temperature and k is a constant.

- For any temperature $T > 0$, function is monotone decreasing function of energy E .
- System more likely to be in a lower energy state than higher one.
 - T large: high and low energy states have roughly same probability
 - T small: low energy states are much more probable


For the vertex cover problem: We could define $E = |S|$.

9

Metropolis Algorithm

Metropolis algorithm.

- Given a fixed temperature T , maintain current state S .
- Randomly perturb current state S to new state $S' \in N(S)$.
- If $E(S') \leq E(S)$, update current state to S'
- Otherwise, update current state to S' with probability $e^{-\Delta E / (kT)}$, where $\Delta E = E(S') - E(S) > 0$.




10

Metropolis Algorithm

Theorem. Let $f_S(t)$ be fraction of first t steps in which simulation is in state S . Then, assuming some technical conditions, with probability 1:

$$\lim_{t \rightarrow \infty} f_S(t) = \frac{1}{Z} e^{-E(S)/(kT)},$$

$$\text{where } Z = \sum_{S \in \text{Universe}} e^{-E(S)/(kT)}.$$



Intuition. Simulation spends roughly the right amount of time in each state, according to Gibbs-Boltzmann equation.

11

Simulated Annealing

Simulated annealing.

- T large \Rightarrow probability of accepting an uphill move is large. State samples explore the space (often computational molasses).
- T small \Rightarrow uphill moves are almost never accepted. State samples remain near local minimum.
- Idea: turn knob to control T .
- Cooling schedule: $T = T(i)$ at iteration i .

Physical analog.

- Take solid and raise it to high temperature, we do not expect it to maintain a nice crystal structure.
- Take a molten solid and freeze it very abruptly, we do not expect to get a perfect crystal either.
- Annealing: cool material gradually from high temperature, allowing it to reach statistical equilibrium at succession of intermediate lower temperatures.


12

Maximum Cut

Maximum cut. Given an undirected graph $G = (V, E)$ with **positive integer** edge weights w_e , find a node partition (A, B) such that the total weight of edges crossing the cut is maximized.

$$w(A, B) := \sum_{u \in A, v \in B} w_{uv}$$

Example problem:



All edge weights 1.
 $A = \text{red}$'s
 $B = \text{blue}$'s
Note: Recolour any single vertex and $w(A, B)$ decreases.
Maximal?

Real applications. **Scheduling problems**, circuit layout, statistical physics.

13

Scheduling Two Machines

Problem Statement. Given jobs $J_n = (p_n, t_n)$, for $n=1, \dots, N$, where t_n is the processing time and p_n is the priority. We wish to assign each job to one of two machines M_1 and M_2 (cannot split a job). Use $m(n) = 1$ or 2 to denote the assignment of job J_n to machine $M_{m(n)}$.

We assume each p_n and t_n are positive integers.

On machine m we sort the jobs, and process them in order of increasing $s_m(n)$. Given this order, the finish time T_n of each job on machine m is:

$$T_n \equiv t_n + \sum_{k: s_m(k) < s_m(n)} t_k.$$

We wish to choose a schedule to minimize the **priority weighted finish time**:

$$F \equiv \sum_{n=1}^N p_n T_n.$$

14

Warm-Up and Review: Scheduling One Machine

If we only had one machine,

$$F_1 \equiv \sum_{n=1}^N p_n T_n = \sum_{n=1}^N \left[p_n t_n + \sum_{k: s(k) < s(n)} p_n t_k \right].$$

A natural greedy rule is to sort the jobs in non-increasing order of priority per unit time, p_n/t_n . Thus, $p_k/t_k > p_n/t_n$ iff $s(k) < s(n)$.

Can prove this job ordering minimizes F_1 (for one machine).

Using this sorting rule we can rewrite F_1 simply as

$$F_1 \equiv \sum_{n=1}^N p_n T_n = \sum_{n=1}^N \left[p_n t_n + \sum_{k < n} w_{kn} \right],$$

where $w_{kn} = \min\{p_k t_n, p_n t_k\}$.

15

Scheduling Two Machines: Graph Formulation

Consider the graph G with vertices $V = \{1, \dots, N\}$ representing each job. The undirected edges E consist of all pairs of vertices with edge weights $w_{kn} = \min\{p_k t_n, p_n t_k\}$.

Suppose (A, B) with $B = V \setminus A$, is a cut corresponding to jobs to be assigned machines M_1 and M_2 respectively. Then summing up the costs for each machine separately we find

$$\begin{aligned} F &\equiv \sum_{n \in A} p_n T_n + \sum_{n \in B} p_n T_n, \\ &= \sum_{n=1}^N p_n t_n + \sum_{k < n, \{k, n\} \subseteq A} w_{kn} + \sum_{k < n, \{k, n\} \subseteq B} w_{kn}, \\ &= \sum_{n=1}^N p_n t_n + \sum_{k < n} w_{kn} - \sum_{j \in A, k \in B} w_{jk}, \\ &= C - w(A, B) \end{aligned}$$

Minimizing objective function F is equivalent to solving max-cut on G !

16

Maximum Cut: Single Flip Neighbourhood

Greedy single-flip. Given a partition (A, B) , move one node from A to B , or one from B to A if it improves the solution.

```
Max-Cut-Local (G, w) {
    Pick a random node partition (A, B)

    while (exists improving node v) {
        if (v is in A) move v to B
        else           move v to A
    }

    return (A, B)
}
```

Local improvement: Suppose $u \in A$, then switch u to B iff

$$\text{Weight of cut edges } (u, v), \text{ if } u \text{ moved to } B. \rightarrow \sum_{v \in A} w_{uv} > \sum_{v \in B} w_{uv} \leftarrow \text{Current weight of cut edges } (u, v), \text{ since } u \in A.$$

17

Greedy Single Flip Runtime

Greedy ascent: On each iteration, **Max-Cut-Local** chooses a single flip which increases $w(A, B)$ (otherwise it stops).

Each iteration can be implemented to run in $O(m+n)$ time, where $|V| = n$, $|E| = m$. (Can simply update the sum of weights for the vertices sharing edges with the one that was switched.)

Lemma: **Max-Cut-Local** runs in $O((n+m)W)$ time.

Pf: Note $w(A, B)$ must increase by at least 1 each iteration (integer edge weights). An upper bound for $w(A, B)$ is $W = \sum_e w_e$, i.e., the sum of all edge weights in G . And a lower bound is 0. Therefore at most W iterations are required. ■

So the runtime is pseudo-polynomial (depends on W not $\log(W)$).

What about the weight of the cut $w(A, B)$ that it produces?

18

Greedy Single Flip Approximation Properties

Theorem. Let (A, B) be a locally optimal partition and let (A^*, B^*) be optimal partition. Then $w(A, B) \geq \frac{1}{2} \sum_e w_e \geq \frac{1}{2} w(A^*, B^*)$.

Pf.

- Local optimality implies that for all $u \in A$: $\sum_{v \in A} w_{uv} \leq \sum_{v \in B} w_{uv}$

weights are nonnegative

↑

↑

negation of local improvement rule

Adding up all these inequalities yields:

$$2 \sum_{\substack{e \in E, \\ e \text{ is } A-A}} w_e \leq \sum_{\substack{u \in A, v \in B}} w_{uv} = w(A, B)$$

↑

negation of local improvement rule

Similarly

$$2 \sum_{\substack{e \in E, \\ e \text{ is } B-B}} w_e \leq \sum_{\substack{u \in A, v \in B}} w_{uv} = w(A, B)$$

↑

negation of local improvement rule

Now,

$$\sum_{e \in E} w_e = \sum_{\substack{e \in A-A \\ \leq \frac{1}{2} w(A, B)}} w_e + \sum_{\substack{e \in B-B \\ \leq \frac{1}{2} w(A, B)}} w_e + \sum_{\substack{e \in A-B \\ w(A, B)}} w_e \leq 2w(A, B) \blacksquare$$

each edge counted once

19

Maximum Cut: Big Improvement Flips

Local search. Has the appropriate a $\frac{1}{2}$ -approximation bound for MAX-CUT, but is not poly-time! (So it is not a $\frac{1}{2}$ -approximation algorithm).

Big-improvement-flip algorithm. Only choose a node which, when flipped, increases the cut value by at least $\frac{2\epsilon}{n} w(A, B)$. That is, suppose $u \in A$, then switch u to B iff

$$\text{Cut edges if } u \text{ moved to } B. \rightarrow \sum_{v \in A} w_{uv} \geq \left(\sum_{v \in B} w_{uv} \right) + \frac{2\epsilon}{n} w(A, B) \leftarrow \text{Current weight of cut edges involving } u.$$

Claim. Upon termination, big-improvement-flip algorithm returns a cut (A, B) with $(2 + \epsilon) w(A, B) \geq w(A^*, B^*)$, i.e. a $1/(2+\epsilon)$ -approximation.

Pf idea. Add $\frac{2\epsilon}{n} w(A, B)$ to each inequality in original proof, find

$$\sum_{e \in E} w_e = \sum_{e \in A-A} w_e + \sum_{e \in B-B} w_e + \sum_{e \in A-B} w_e \leq (2 + \epsilon) w(A, B)$$

20

Maximum Cut: Big Improvement Flips

Claim. Big-improvement-flip algorithm terminates after $O(\varepsilon^{-1} n \log W)$ flips, where $W = \sum_e w_e$.

That is, big-improvement-flip is a $1/(2+\varepsilon)$ -approximation algorithm.

Pf:

- Each flip improves cut value by at least a factor of $(1 + \varepsilon/n)$.
- After n/ε iterations the cut value improves by a factor of $m = (1 + \varepsilon/n)^{n/\varepsilon}$.
- For $x \geq 1$, $f(x) = (1+1/x)^x \geq 2$. [Blackboard]
- So for $n/\varepsilon \geq 1$, the cut value is at least doubled in n/ε iterations.
- But integer cut value can be doubled at most $O(\log(W))$ times.
- Total number of iterations is therefore $O((n/\varepsilon)\log(W))$.

21

Neighbor Relations for Max Cut

1-flip neighborhood. (A, B) and (A', B') differ in **exactly one** node.

k-flip neighborhood. (A, B) and (A', B') differ in **at most k** nodes.

- $\Theta(n^k)$ neighbors.

KL-neighborhood. [Kernighan-Lin, 1970]

cut value of (A_1, B_1) may be
worse than (A, B)

- To form neighborhood of (A, B) :
 - Iteration 1: flip node from (A, B) that results in best cut value (A_1, B_1) , and mark that node.
 - Iteration i: flip node from (A_{i-1}, B_{i-1}) that results in best cut value (A_i, B_i) among all nodes not yet marked.
- Neighborhood of $(A, B) = (A_1, B_1), \dots, (A_{n-1}, B_{n-1})$.
- Neighborhood includes some very long sequences of flips, but without the computational overhead of a k-flip neighborhood.
- Practice: powerful and useful framework.
- Theory: explain and understand its success in practice.

22

Maximum Cut: Context

Theorem. [Shani-Gonzales, 1976] There exists a $\frac{1}{2}$ -approximation algorithm (polytime by defn) for MAX-CUT.

Theorem. [Goemans-Williamson, 1995] There exists an 0.878567-approximation algorithm for MAX-CUT.

$$\min_{0 \leq \theta \leq \pi} \frac{2}{\pi} \frac{\theta}{1 - \cos \theta}$$

The G-W approach uses semi-definite programming (SDP) and provides a useful rounding heuristic in practice.


Theorem. [Håstad, 2001] Unless $P = NP$, no 16/17 approximation algorithm for MAX-CUT.

0.941176

23

Example: Local Search in Shape from Texture


Humans can perceive shape from image texture. Can a machine?



24

Shape From Projected Squares

Depth-flip ambiguity:




Given the image of a small square, there are two possible 3D orientations.

Given a mathematical model for image projection, we can solve for these orientations. We refer to them by $d = \pm 1$.

25

Merging Local Depth Flips

Every image parallelogram has a depth-flip ambiguity $d_k = \pm 1$.



We wish to find a set of flips that can be fit with a smooth surface.

26